Gradient-probe trajectories — classic routing, no coherence (1 seed/env; median + IQR over probe positions)

forget share of gradient norm

gradient preference Λ

co-implement (+) vs counteract (−): solid=hack onsets, dashed=clean

receptivity: solid=hack onsets, dashed=clean

saturation context

behavior emergence (run's own eval)

addition_v2 s1 (conditional)

addition_v2 s3 (localized)

cities_qa

persona_qa

repeat_extra

sorting_copy

topic_contains

MEAN

ρ_F / (ρ_R + ρ_F)

Λ = ρ_F/ρ_R

signed alignment ŝ

ρ (norm. grad)

per-token CE loss

hack frequency (both)

hack frequency

Λ

ŝ

CE loss

hack onset (blind-spot)

hack onset (detected)

clean tokens

ρ_R hack-onset

ρ_F hack-onset

ρ_R clean

ρ_F clean

ŝ_R hack-onset

ŝ_F hack-onset

ŝ_R clean

ŝ_F clean

hack onset

clean

training step

fraction of training